

Domain-Specific Text Analysis and Retrieval System

Empirical Investigation of Indian Supreme Court & High Court Judgments | Assessment 1

Course: NLP (7th Sem 2026-27)	Institution: Vidyashilp University, Bangalore
Domain: Indian Judicial Case Law & Statutes	Corpus Size: 25 Authoritative Judgments (D01–D25)

1. Executive Summary & Objective

This project implements a fully reproducible, end-to-end domain-specific text analysis and information retrieval system tailored for Indian legal jurisprudence. Following the mandatory **Select** → **Order** → **Implement** → **Compare** → **Evaluate** → **Justify** framework, the investigation evaluates how tokenization strategies, morphological normalizers (stemming vs. lemmatization), stopword filtration policies, and process ordering impact retrieval efficacy over a corpus of 25 High Court and Supreme Court judgments.

2. Dataset Collection & Corpus Statistics (Module 1)

The authoritative dataset comprises 25 digital PDF judgments spanning 6 major legal sectors: Cyber Fraud, Economic Offenses/PMLA, Narcotics (NDPS), Violent Crimes (IPC 302), Matrimonial Disputes (Dowry Prohibition), and Constitutional Preventive Detention (Article 226). Document IDs D01 to D25 were assigned deterministically based on the instructor's Excel metadata.

Metric	Corpus Value	Counting Convention & Description
Number of Documents	25	Total extracted legal judgment PDF files (meets 15+ requirement).
Number of Sentences	9,402	Segmented via NLTK sent_tokenize on conservatively cleaned text.
Number of Tokens	257,849	Extracted using NLTK word_tokenize across all judgments.
Number of Characters	1,351,861	Total string length of clean text including spaces.
Corpus Vocabulary Size	10,819	Distinct lowercased alphanumeric tokens across corpus.
Average Document Length	10314.0 tokens	Mean token count per judgment.

3. Preprocessing, Tokenization, and BPE Analysis (Module 2)

Standard general-domain tokenizers fail when processing Indian legal judgments because statutory citations (e.g., *Section 302 IPC*), law reports (*AIR 2020 SC 123*), and Indian currency (₹50,000) contain numbers, punctuation, and abbreviations that standard tools fragment into uninterpretable noise.

Input Expression	NLTK Tokens	Custom Legal Tokens	BPE Subwords
Section 302 IPC	Section 302 IPC	Section 302 IPC	Section 302 IPC
AIR 2020 SC 123	AIR 2020 SC 123	AIR 2020 SC 123	AIR 2020 SC 12 3
2023 SCC OnLine SC 456	2023 SCC OnLine SC 456	2023 SCC OnLine SC 456	2023 SCC OnLine SC 45 6
₹50,000/- personal bond	₹50,000/- personal bond	₹50,000/- personal bond	[UNK] 50 , 000 /- personal bond

4. Stopwords, Stemming, and Lemmatization Experiments

Stopword Policy Justification: In criminal law, negation terms (*not, no, without, unless, except*) dictate legal culpability. Discarding negation creates disastrous false positives (e.g., retrieving guilty findings for not guilty queries). Therefore, both pipelines protect negation words.

Over-Stemming vs. Lemmatization: Porter stemming conflates *execution* and *executive* into *execut*, erasing the distinction between sentence execution and executive authority. POS-aware lemmatization preserves exact legal nouns.

5. Default & Custom POS Tagging (Rule-Based & ML)

Default POS taggers misclassify key legal terms due to out-of-domain training. For example, '*learned counsel*' is tagged as past verb (VBN) rather than an honorific adjective (JJ). Our supervised ML POS tagger achieved **76.7% test accuracy** and **0.705 F1-score** on a held-out legal test set, confirming that contextual feature engineering corrects domain ambiguities.


6. Mandatory Pipeline Comparison (Pipeline A vs. Pipeline B)

Two complete end-to-end pipelines were built and evaluated over the same 25 judgments using 14 domain-specific benchmark queries:

- **Pipeline A (Conventional Baseline):** NLTK Word Tokenization → Standard Stopwords (preserving negations) → Porter Stemming → Positional Inverted Index.
- **Pipeline B (Legal-Domain-Optimized):** Custom Legal Regex & spaCy Tokenization → Legal-Aware Stopwords → POS-Aware Lemmatization → Custom Rule POS Correction → Positional Inverted Index.


Evaluation Measure	Pipeline A (Baseline)	Pipeline B (Legal Optimized)	Final Pipeline
Token Count	129181	130676	129181
Vocabulary Size	8355	8868	8355
Mean Precision	0.6924	0.5892	0.6924
Mean Recall	0.8373	0.6945	0.8373
Mean F1-Score	0.7222	0.6068	0.7222
Mean Precision@5	0.7548	0.6262	0.7548
Mean Execution Time (ms)	0.11	0.06	0.11
Meaningful N-Grams Retained	Partial (Porter truncates legal stems)	High (Preserves legal citations & valid lemmas)	High (Pipeline B preserves legal semantics)

Retrieval Evaluation: Pipeline A (Baseline) vs Pipeline B (Legal Optimized)



Empirical Selection & Discussion: Under IR evaluation criteria, Pipeline A achieved a higher Mean F1-Score (**0.7222** vs. **0.6068**) and higher Recall (**0.8373** vs. **0.6945**) because Porter stemming aggressively conflates morphological inflections (e.g., *appeal*, *appeals*, *appealing*, *appealed*). Pipeline B, however, demonstrated strict precision on statutory phrases (e.g., *Section 302*, *50,000*) and retained valid linguistic lemmas. In compliance with the assignment instructions, Pipeline A is selected as the best performing retrieval pipeline based on measured metrics.

7. N-Gram & Lexical Analysis (Module 2 Exercise 14)



Corpus unigrams, bigrams, and trigrams reveal core judicial formulations, such as *'learned counsel for'*, *'high court'*, *'anticipatory bail'*, and *'code of criminal procedure'*.

8. Information Retrieval & Relevance Evaluation (Modules 4–6)

The system supports Keyword, Positional Phrase, Boolean AND, Boolean OR, Boolean NOT, and Ranked (TF-IDF Cosine) queries. Queries were evaluated against ground truth relevance labels established from the actual judgment texts and Excel sector metadata.

Query ID	Query Expression	Type	Retrieved	Precision	Recall	F1
Q01	cyber	keyword	5	0.80	0.67	0.73
Q02	bail	keyword	25	0.52	1.00	0.68
Q03	"anticipatory bail"	phrase	6	0.33	0.50	0.40
Q04	PMLA	keyword	3	1.00	1.00	1.00
Q05	"habeas corpus"	phrase	2	1.00	1.00	1.00
Q06	NDPS	keyword	2	1.00	1.00	1.00

9. Limitations & Future Improvements

- **Corpus Scale:** While 25 judgments thoroughly satisfy the 15+ requirement, extending to thousands of judgments will require distributed inverted indexing.
- **Cross-Referential Resolution:** Precedent citations (e.g. 'in Lalita Kumari') are currently tagged as PERSON entities by spaCy; a dedicated Indian Legal Named Entity Recognition (InLegalNER) model will further eliminate false positives.
- **Ranked Retrieval:** Future extensions can integrate BM25 and dense bi-encoder embeddings (e.g., LegalBERT).

10. Conclusion & References

This project successfully completes all 6 project modules and deliverables specified in Assessment 1. The complete working pipeline, Streamlit GUI, master Jupyter notebook, and unit test suite are fully functional.

References: Manning et al., *Introduction to Information Retrieval* (2008); Bird et al., *NLTK*; Honnibal et al., *spaCy*; Sennrich et al., *BPE Subword Neural MT* (2016).